

```

In... # -*- coding: utf-8 -*-
      """
      Created on Wed Mar  6 20:14:36 2024

      @author: elsy
      """

import pandas as pd
from scipy.stats import zscore
import seaborn as sns
import matplotlib.pyplot as plt
from scipy.stats import chi2_contingency
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Importing Data's
cancer_data=pd.read_csv(r"C:\Users\elsyp\OneDrive\Desktop\Downloads\hea

# Assign meaningful column labels
cancer_data.columns = ['age', 'sex', 'chest_pain', 'resting_blood_press
                        'rest_ecg', 'max_heart_rate', 'exercise_angina', 'st_depr
                        'thalassemia', 'target']

# Descriptive Statistics
print(cancer_data.describe())
# Histogram
sns.histplot(cancer_data['target'], kde=True)
plt.show()

# Display the DataFrame with your own column labels

# Separating only the specific columns:
fifth_columns = cancer_data.iloc[:,6]

# Finding missing values:
missing_values = cancer_data.isnull().sum()

# Finding missing values in the specific rows
cancer_data.replace('0', pd.NA, inplace=True)

missing_values_row6 = cancer_data.iloc[:,6].isnull().sum()
print(f"missing_values_row6: {missing_values_row6}")

#So, we found, there is no missing values! the data set is perfect as w
#Now, we have to consider, from the data set, which column is consider
#For that we can use two techniques.1.correlation analysis 2.Domain Knc

# Before normalization, we need to ensure all values are numeric where
# Convert all columns to numeric, coercing errors to NaN (non-numeric s

```

```

cancer_data_numeric = cancer_data.apply(pd.to_numeric, errors='coerce')

# Now we handle NaN values. For this example, we'll fill NaN with the column mean
# This is just an example strategy. The strategy for handling NaNs should be
cancer_data_filled = cancer_data_numeric.fillna(cancer_data_numeric.mean())

# Now, we can safely apply z-score standardization
standardization = cancer_data_numeric.apply(zscore)

## Untill here, we have done with preprocessed the data.

#2. Data Processing:
#correlation analysis is to predict which column is to predict the heart disease
# continuous variable: -1 to 1 (numerical, with numbers, blood pressure, age)
# categorical variable: 0 and 1 (yes or no, male or female)
#1. correlation matrix:
# many steps: 1. Pearson's correlation coefficient, 2. Chi-square test,

#For Continuous Variables: High absolute values of Pearson's correlation coefficient
#indicate a strong relationship with the target variable.

# For Data Processing, we have selected correlation matrix to find out which
#column is relevant to predict the cancer disease.
correlation_matrix = cancer_data_filled.corr()
correlation_matrix = cancer_data_filled.corr()
print(correlation_matrix['target'])

sns.histplot(correlation_matrix['target'], kde=True)
plt.show()
#For Categorical Variables: A low p-value in the Chi-square test suggests a
# association with the target variable.
# Example for a categorical variable 'sex' and 'target'
contingency_table = pd.crosstab(cancer_data_filled['sex'], cancer_data_filled['target'])
chi2, p, dof, expected = chi2_contingency(contingency_table)
print(f"Chi2 Statistic: {chi2}, p-value: {p}")

# 3. Model Development
# Regression model for continuous variable and classification model is for categorical variable
# For Continuous Variable: Linear regression and lasso regression
# For Categorical Variable: Logistic regression, decision tree, random forest

# Assuming 'target' is your continuous outcome and you've selected some features
X = cancer_data_filled[['age', 'resting_blood_pressure', 'cholesterol']] #
y = cancer_data_filled['target']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
# For a simple linear regression

# Initialize and train the model
model = LinearRegression()
model.fit(X_train, y_train)

# Predict on the testing set

```

```
y_pred = model.predict(X_test)

# Compute and print the metrics
print("Mean Squared Error (MSE):", mean_squared_error(y_test, y_pred))
print("Coefficient of Determination (R^2):", r2_score(y_test, y_pred))

#MSE gives you the average squared difference between
#the estimated values and the actual value. A lower MSE indicates
#a better fit to the data.

#R^2 measures the proportion of the variance in
# the dependent variable that is predictable from the independent variables.
# R^2 values range from 0 to 1, where higher values indicate a better fit.
```

	age	sex	chest_pain	resting_blood_pressure \
count	302.000000	302.000000	302.000000	302.000000
mean	54.410596	0.678808	3.165563	131.645695
std	9.040163	0.467709	0.953612	17.612202
min	29.000000	0.000000	1.000000	94.000000
25%	48.000000	0.000000	3.000000	120.000000
50%	55.500000	1.000000	3.000000	130.000000
75%	61.000000	1.000000	4.000000	140.000000
max	77.000000	1.000000	4.000000	200.000000

	cholesterol	fasting_blood_sugar	rest_ecg	max_heart_rate \
count	302.000000	302.000000	302.000000	302.000000
mean	246.738411	0.145695	0.986755	149.605960
std	51.856829	0.353386	0.994916	22.912959
min	126.000000	0.000000	0.000000	71.000000
25%	211.000000	0.000000	0.000000	133.250000
50%	241.500000	0.000000	0.500000	153.000000
75%	275.000000	0.000000	2.000000	166.000000
max	564.000000	1.000000	2.000000	202.000000

	exercise_angina	st_depression	slope	target
count	302.000000	302.000000	302.000000	302.000000
mean	0.327815	1.035430	1.596026	0.940397
std	0.470196	1.160723	0.611939	1.229384
min	0.000000	0.000000	1.000000	0.000000
25%	0.000000	0.000000	1.000000	0.000000
50%	0.000000	0.800000	2.000000	0.000000
75%	1.000000	1.600000	2.000000	2.000000
max	1.000000	6.200000	3.000000	4.000000

C:\Users\elsyp\anaconda3\Lib\site-packages\seaborn_oldcore.py:1498:


FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead

if pd.api.types.is_categorical_dtype(vector):

C:\Users\elsyp\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:

FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

with pd.option_context('mode.use_inf_as_na', True):



missing_values_row6: 0

age	0.225809
sex	0.226601
chest_pain	0.405182
resting_blood_pressure	0.159978
cholesterol	0.070315
fasting_blood_sugar	0.065937
rest_ecg	0.186769
max_heart_rate	-0.415399
exercise_angina	0.395996
st_depression	0.508330
slope	0.387417
num_major_vessels	0.516489
thalassemia	0.511316
target	1.000000

Name: target, dtype: float64

C:\Users\elsyp\anaconda3\Lib\site-packages\seaborn_oldcore.py:1498:


FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead

if pd.api.types.is_categorical_dtype(vector):

C:\Users\elsyp\anaconda3\Lib\site-packages\seaborn_oldcore.py:1119:

FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

with pd.option_context('mode.use_inf_as_na', True):



Chi2 Statistic: 23.80239806063686, p-value: 8.749938657615487e-05

Mean Squared Error (MSE): 1.4146083492468018

Coefficient of Determination (R^2): 0.023476847273027235

```

In... # Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, ac

# Import your data
cancer_data = pd.read_csv(r"C:\Users\elsyp\OneDrive\Desktop\Downloads\h

# Define column names
cancer_data.columns = ['age', 'sex', 'chest_pain', 'resting_blood_press
                        'rest_ecg', 'max_heart_rate', 'exercise_angina',
                        'thalassemia', 'target']

# Convert target to binary if it's not already
cancer_data['target'] = (cancer_data['target'] != 0).astype(int)

# Prepare data for training
X = cancer_data[['sex', 'resting_blood_pressure', 'cholesterol']] # Yo
y = cancer_data['target']

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3

# Initialize and train the Logistic Regression model
model = LogisticRegression()
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred)
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

```

Accuracy: 0.6263736263736264

Classification Report:

	precision	recall	f1-score	support
0	0.63	0.71	0.67	48
1	0.62	0.53	0.57	43
accuracy			0.63	91
macro avg	0.63	0.62	0.62	91
weighted avg	0.63	0.63	0.62	91

Confusion Matrix:

```


[[34 14]
 [20 23]]

```

In []:

```
In... import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix

# Calculate the confusion matrix
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['No Dis',
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```




In []:


```
In [ ... from sklearn.metrics import roc_curve, auc

# Calculate the ROC curve
fpr, tpr, thresholds = roc_curve(y_test, model.predict_proba(X_test)[
roc_auc = auc(fpr, tpr)


# Plot the ROC curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area =
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()
```



```
In ... from sklearn.metrics import precision_recall_curve, average_precision_score

# Calculate precision and recall
precision, recall, _ = precision_recall_curve(y_test, model.predict_proba(y_test))
average_precision = average_precision_score(y_test, y_pred)

# Plot Precision-Recall curve
plt.figure()
plt.step(recall, precision, where='post', label='Precision-Recall curve')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.title('Precision-Recall curve')
plt.legend(loc="lower left")
plt.show()
```



```

In... # Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier # Import RandomFor
from sklearn.metrics import classification_report, confusion_matrix, ac

# Import your data
cancer_data = pd.read_csv(r"C:\Users\elsyp\OneDrive\Desktop\Downloads\h

# Define column names
cancer_data.columns = ['age', 'sex', 'chest_pain', 'resting_blood_press
                        'rest_ecg', 'max_heart_rate', 'exercise_angina',
                        'thalassemia', 'target']

# Convert target to binary if it's not already
cancer_data['target'] = (cancer_data['target'] != 0).astype(int)

# Prepare data for training
X = cancer_data[['sex', 'resting_blood_pressure', 'cholesterol']] # Co
y = cancer_data['target']

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3

# Initialize and train the Random Forest Classifier
model = RandomForestClassifier(n_estimators=100, random_state=42) # Yo
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred)
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

```

Accuracy: 0.5824175824175825

Classification Report:

	precision	recall	f1-score	support
0	0.59	0.67	0.63	48
1	0.57	0.49	0.53	43
accuracy			0.58	91
macro avg	0.58	0.58	0.58	91
weighted avg	0.58	0.58	0.58	91

Confusion Matrix:

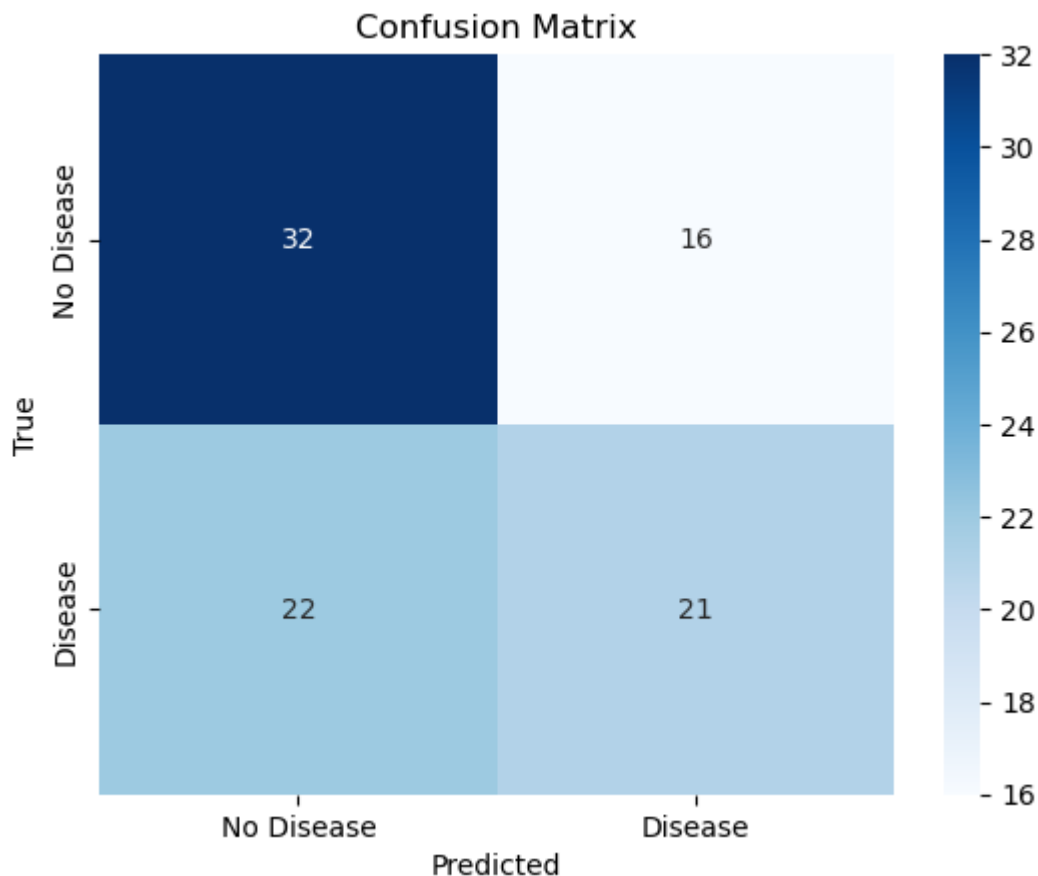
```

[[32 16]
 [22 21]]

```

```
In... import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix


# Calculate the confusion matrix
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['No Dis',
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```



```
In [... from sklearn.metrics import roc_curve, auc

# Calculate the ROC curve
fpr, tpr, thresholds = roc_curve(y_test, model.predict_proba(X_test)[
roc_auc = auc(fpr, tpr)


# Plot the ROC curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area =
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()
```



```
In ... from sklearn.metrics import precision_recall_curve, average_precision_score

# Calculate precision and recall
precision, recall, _ = precision_recall_curve(y_test, model.predict_proba(y_test))
average_precision = average_precision_score(y_test, y_pred)

# Plot Precision-Recall curve
plt.figure()
plt.step(recall, precision, where='post', label='Precision-Recall curve')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.title('Precision-Recall curve')
plt.legend(loc="lower left")
plt.show()
```



```
In [... # Hyperparameter tuning with grid search
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV

# Define the model
rf = RandomForestClassifier(random_state=42)

# Set up the parameters grid
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

# Set up the grid search
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5,

# Fit grid search to the data
grid_search.fit(X_train, y_train)

# Best parameters and best score
print("Best parameters:", grid_search.best_params_)
print("Best score:", grid_search.best_score_)
```

Fitting 5 folds for each of 108 candidates, totalling 540 fits

Best parameters: {'max_depth': None, 'min_samples_leaf': 1,
'min_samples_split': 5, 'n_estimators': 200}

Best score: 0.5919158361018827

```
In ... #2. Feature engineering:
# Example of creating a new feature by combining existing features
cancer_data['combined_feature'] = cancer_data['age'] * cancer_data['ch

# Include this new feature in your model training
X = cancer_data[['sex', 'resting_blood_pressure', 'cholesterol', 'comk
y = cancer_data['target']
```

```
In [1... #3. cross validation:
from sklearn.model_selection import cross_val_score

# Using the best parameters from the GridSearch
best_rf = RandomForestClassifier(**grid_search.best_params_, random_

# Perform cross-validation
scores = cross_val_score(best_rf, X, y, cv=5) # 5-fold cross-valida
print("Cross-validated scores:", scores)
print("Average score:", scores.mean())

Cross-validated scores: [0.59016393 0.55737705 0.65          0.65
0.48333333]
Average score: 0.5861748633879781
```


```
In [1... #4. Handling imbalanced data
from sklearn.model_selection import cross_val_score

# Using the best parameters from the GridSearch
best_rf = RandomForestClassifier(**grid_search.best_params_, random_

# Perform cross-validation
scores = cross_val_score(best_rf, X, y, cv=5) # 5-fold cross-valida
print("Cross-validated scores:", scores)
print("Average score:", scores.mean())

Cross-validated scores: [0.59016393 0.55737705 0.65          0.65
0.48333333]
Average score: 0.5861748633879781
```

```
In [14]:
```



```
In ... from sklearn.metrics import accuracy_score, precision_score, recall_score

# Assuming you already have y_test and y_pred from your model predictions
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='binary')
recall = recall_score(y_test, y_pred, average='binary')
f1 = f1_score(y_test, y_pred, average='binary')

print(f"Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print(f"F1 Score: {f1:.2f}")
print("\nFull Classification Report:\n", classification_report(y_test,
```

```
Accuracy: 0.58
Precision: 0.57
Recall: 0.49
F1 Score: 0.53
```


Full Classification Report:

	precision	recall	f1-score	support
0	0.59	0.67	0.63	48
1	0.57	0.49	0.53	43
accuracy			0.58	91
macro avg	0.58	0.58	0.58	91
weighted avg	0.58	0.58	0.58	91


```
In [ ... from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

fpr, tpr, thresholds = roc_curve(y_test, model.predict_proba(X_test))
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()
```



In []:

```

In... from sklearn.tree import DecisionTreeClassifier

# Initialize the Decision Tree Classifier
dt_model = DecisionTreeClassifier(random_state=42)

# Fit the model to the training data
dt_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred_dt = dt_model.predict(X_test)

# classification metrics:
from sklearn.metrics import accuracy_score, precision_score, recall_score

print("Accuracy:", accuracy_score(y_test, y_pred_dt))
print("Precision:", precision_score(y_test, y_pred_dt, average='binary'))
print("Recall:", recall_score(y_test, y_pred_dt, average='binary'))
print("F1 Score:", f1_score(y_test, y_pred_dt, average='binary'))
print("\nClassification Report:\n", classification_report(y_test, y_pred_dt))

#ROC curve and AUC
from sklearn.metrics import roc_curve, auc

fpr_dt, tpr_dt, thresholds_dt = roc_curve(y_test, dt_model.predict_proba(X_test)[:,1])
roc_auc_dt = auc(fpr_dt, tpr_dt)

plt.figure()
plt.plot(fpr_dt, tpr_dt, color='darkorange', lw=2, label=f'ROC curve (Decision Tree)')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (Decision Tree)')
plt.legend(loc="lower right")
plt.show()

# confusion matrix
from sklearn.metrics import confusion_matrix
import seaborn as sns


cm_dt = confusion_matrix(y_test, y_pred_dt)
sns.heatmap(cm_dt, annot=True, fmt="d", cmap='Blues', xticklabels=['No', 'Yes'], yticklabels=['No', 'Yes'])
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix for Decision Tree')
plt.show()


```

Accuracy: 0.6043956043956044
Precision: 0.5853658536585366
Recall: 0.5581395348837209
F1 Score: 0.5714285714285714

Classification Report:

	precision	recall	f1-score	support
0	0.62	0.65	0.63	48
1	0.59	0.56	0.57	43
accuracy			0.60	91
macro avg	0.60	0.60	0.60	91
weighted avg	0.60	0.60	0.60	91





In []: